

Floating-point arithmetic and error analysis (AFAE)

Fast verification methods for linear systems – Part II

Stef Graillat

LIP6/PEQUAN – Sorbonne University

Lecture Master 2 SFPN – MAIN5



1 Iterative refinement

- Error analysis
- Computation of a verified solution
- Numerical experiments

2 Inversion of ill-conditioned matrices

- Statement of the problem
- Rump's algorithm

3 Conclusion

Problem statement

Let $A \in \mathbb{F}^{n \times n}$ be nonsingular, $b \in \mathbb{F}^n$, and $\hat{x} \in \mathbb{F}^n$ an approximate solution of the linear system $Ax = b$. We have already seen four “fast” to compute an upper bound for $\|\hat{x} - x\|_\infty$.

How can we increase the accuracy of \hat{x} if this one is not sufficient?

One possibility is to redo all the computations using extended precision.

Another possibility is to use **iterative refinement** to increase the accuracy of the computed solution.

Iterative refinement

- The iterative refinement is a technique to increase the accuracy of the computed solution \hat{x} of a linear system $Ax = b$:
 - 1 $\hat{x}_i \leftarrow$ computed solution of $Ax = b$
 - 2 $\hat{r}_i \leftarrow$ computed residual $b - A\hat{x}_i$
 - 3 $\hat{c}_i \leftarrow$ computed solution of $Ac_i = \hat{r}_i$
 - 4 $\hat{x}_{i+1} \leftarrow \text{fl}(\hat{x}_i + \hat{c}_i)$
 - 5 go to step 2 if the stopping criterion is not satisfied
- We generally distinguish 2 cases:
 - 1 either the residual is computed with the working precision: this makes it possible to increase the backward error of the computed solution [Hig02, Thm 12.1 and 12.2],
 - 2 or the residual is computed with twice the working precision: this makes it possible to increase the forward error [Hig02, Thm 12.1].

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

Conditioning

We introduce the following componentwise condition number:

$$\text{cond}_{E,f}(A, x) := \lim_{\varepsilon \rightarrow 0} \sup_{\substack{|\Delta A| \leq \varepsilon |E| \\ |\Delta b| \leq \varepsilon |f|}} \left\{ \frac{\|\hat{x} - x\|_\infty}{\varepsilon \|x\|_\infty}, (A + \Delta A) \hat{x} = b + \Delta b \right\}.$$

If we take $E = |A|$ and $f = |b|$, we denote $\text{cond}(A, x) := \text{cond}_{A,b}(A, x)$, and we have

$$\text{cond}(A, x) = \frac{\| |A^{-1}| |A| |x| \|_\infty}{\|x\|_\infty}.$$

If we take $E = |A|$ and $f = 0$, we denote $\text{cond}(A) := \text{cond}_{A,0}(A, x)$, and we have

$$\text{cond}(A) = \| |A^{-1}| |A| \|_\infty.$$

Let us remark that $\text{cond}(A, x) \leq \text{cond}(A) \leq \kappa_\infty(A)$.

Error bound

We have already proved the following result:

Theorem 1 (Thm 9.4, p. 164, in [Hig02])

Let $A \in \mathbb{F}^{n \times n}$ and suppose GE produces computed LU factors $A \approx \hat{L}\hat{U}$, and a computed solution \hat{x} to $Ax = b$. Then

$$(A + \Delta A)\hat{x} = b, \quad |\Delta A| \leq \gamma_{3n} |\hat{L}| |\hat{U}|.$$

This is not entirely satisfactory: in the normwise sense, even with GE, “the solution is stable only if the growth factor is small”.

In this case, we can show that

$$\frac{\|\hat{x} - x\|_{\infty}}{\|x\|_{\infty}} \leq \gamma_{3n} \frac{\|A^{-1}\| \|\hat{L}\| \|\hat{U}\| \|\hat{x}\|_{\infty}}{\|x\|_{\infty}},$$

where the second factor in the RHS is to be compared with $\text{cond}(A, x)$.

Refinement with fixed precision u

We assume that for all matrix $A \in \mathbb{F}^{n \times n}$ and for all vector $b \in \mathbb{F}^n$, the method used for solving $Ax = b$ returns an approximate solution satisfying

$$(A + \Delta A) \hat{x} = b, \quad \text{with} \quad |\Delta A| \leq uW(A, n).$$

It is what we obtain with GE; in this case, indeed, \hat{x} satisfies

$$(A + \Delta A) \hat{x} = b, \quad \text{with} \quad |\Delta A| \leq \gamma_{3n} |\hat{L}| |\hat{U}|.$$

Theorem 2 (Thm 12.2, p. 234, in [Hig02])

Let iterative refinement in fixed precision be applied to the nonsingular linear system $Ax = b$ of order n . Let $\eta = u \| |A^{-1}| (|A| + W(A, n)) \|_{\infty}$. Then, provided η is sufficiently less than 1, iterative refinement reduces the forward error by a factor approximately η at each stage, until

$$\frac{\| \hat{x}_k - x \|_{\infty}}{\| x \|_{\infty}} \leq 2nu \operatorname{cond}(A, x) + O(u^2).$$

- If GE is used, $\eta = u \| |A^{-1}| (|A| + 3n | \hat{L} | | \hat{U} |) \|_{\infty} + O(u^2)$:
if $| \hat{L} | | \hat{U} | \approx |A|$, then $\eta \approx 3nu \operatorname{cond}(A)$.
- An upper bound in $u \operatorname{cond}(A, x)$ is the best we can expect with working in precision u .
- The theorem gives a sharper bound than the initial one

$$\frac{\| \hat{x} - x \|_{\infty}}{\| x \|_{\infty}} \leq \gamma_{3n} \frac{\| |A^{-1}| | \hat{L} | | \hat{U} | \hat{x} \|_{\infty}}{\| x \|_{\infty}}.$$

Refinement with mixed precision

Theorem 3 ([Hig02, § 12.1])

Let iterative refinement be applied to the nonsingular linear system $Ax = b$ with residuals computed in double the working precision. Let $\eta = u \| |A^{-1}|(|A| + W(A, n)) \|_{\infty}$. Then, provided η is sufficiently less than 1, iterative refinement reduces the forward error by a factor approximately η at each stage until

$$\frac{\| \hat{x}_k - x \|_{\infty}}{\| x \|_{\infty}} \approx u.$$

- This time, we obtain a forward error of the order of u if the condition number is sufficiently small
- By allowing to double the precision used to compute the residuals, we can obtain a forward error that is smaller than $u \operatorname{cond}(A, x)$.

Outline

- 1 **Iterative refinement**
 - Error analysis
 - **Computation of a verified solution**
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

We will use iterative refinement into those certified algorithm:

- certifLSV1 : $6 \cdot n^3 + O(n^2)$;
- certifLSV4 : $4/3 \cdot n^3 + O(n^2)$.

The algorithms will have the following form:

- 1 $\hat{x}_i \leftarrow$ computed solution of $Ax = b$
- 2 $\hat{r}_i \leftarrow$ computed residual $b - A\hat{x}_i$ { with double the working precision }
- 3 $\hat{c}_i \leftarrow$ computed solution of $Ac_i = \hat{r}_i$
- 4 $\hat{x}_{i+1} \leftarrow \text{fl}(\hat{x}_i + \hat{c}_i)$
- 5 Computation of a verified upper bound δ for $\|\hat{x}_{i+1} - x\|_\infty$
- 6 Go to step 2 if the stopping criterion is not valid

Choosing a stopping criterion

- Assume that we know an upper bound $\bar{\delta}$ such that $\|\mathbf{x} - \hat{\mathbf{x}}\|_{\infty} \leq \bar{\delta}$.

We have $\bar{\delta} \geq \|\hat{\mathbf{x}}\|_{\infty} - \|\mathbf{x}\|_{\infty}$, so $\|\mathbf{x}\|_{\infty} \geq \|\hat{\mathbf{x}}\|_{\infty} - \bar{\delta}$.

Hence, **assuming that $\bar{\delta} < \|\hat{\mathbf{x}}\|_{\infty}$** , we obtain:

$$\frac{\|\mathbf{x} - \hat{\mathbf{x}}\|_{\infty}}{\|\mathbf{x}\|_{\infty}} \leq \frac{\bar{\delta}}{\|\hat{\mathbf{x}}\|_{\infty} - \bar{\delta}}.$$

From algorithm `bndDelta1`, we can deduce an algorithm `bndEpsilon`, to upper bound the relative error

$$\|\mathbf{x} - \hat{\mathbf{x}}\|_{\infty} / \|\mathbf{x}\|_{\infty}.$$

- We choose to stop the iterative refinement when:
 - either $\|\mathbf{x} - \hat{\mathbf{x}}\|_{\infty} / \|\mathbf{x}\|_{\infty} \leq \tau$, where τ is the tolerance;
 - or 3 steps of iterative refinement have already been done.

Verified solution of linear system with iterative refinement

```
function [ $\hat{\mathbf{x}}$ ,  $\bar{\epsilon}$ ] = certifLSV1rafit(A, b,  $\tau$ )  
  [L, U, P] = fl (xGETRF(A)) {  $2/3 \cdot n^3 + O(n^2)$  }  
   $\hat{\mathbf{x}}$  = fl (xGETRS(P, L, U, b)) {  $O(n^2)$  }  
  R = fl (xGETRI(P, L, U)) {  $4/3 \cdot n^3 + O(n^2)$  }  
   $\bar{\alpha}$  = bndAlpha1(A, R) {  $4 \cdot n^3 + O(n^2)$  }  
  if  $\bar{\alpha} \geq 1$  then error('Certification failed')  
  k = 0; while true do  
     $\langle m_{\text{res}}, r_{\text{res}} \rangle$  = resLinSys2(A, b,  $\hat{\mathbf{x}}$ ) {  $O(n^2)$  }  
     $\bar{\epsilon}$  = bndEpsilon(A, b,  $\hat{\mathbf{x}}$ , R,  $\langle m_{\text{res}}, r_{\text{res}} \rangle$ ) {  $O(n^2)$  }  
    if  $\bar{\epsilon} \leq \tau$  then return  
    if  $k \geq 3$  then error('Convergence failed')  
     $\hat{\mathbf{c}}$  = fl (xGETRS(P, L, U,  $m_{\text{res}}$ )) {  $O(n^2)$  }  
     $\hat{\mathbf{x}}$  = fl( $\hat{\mathbf{x}}$  +  $\hat{\mathbf{c}}$ )  
    k = k + 1  
  done
```

Cost of the algorithm: $6 \cdot n^3 + O(n^2)$.

Outline

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

All the numerical experiments were done the the same environment as previously:

- MATLAB, with INTLAB toolbox;
- IEEE-754 double precision;
- the ill-conditioned linear systems are generated as follows:

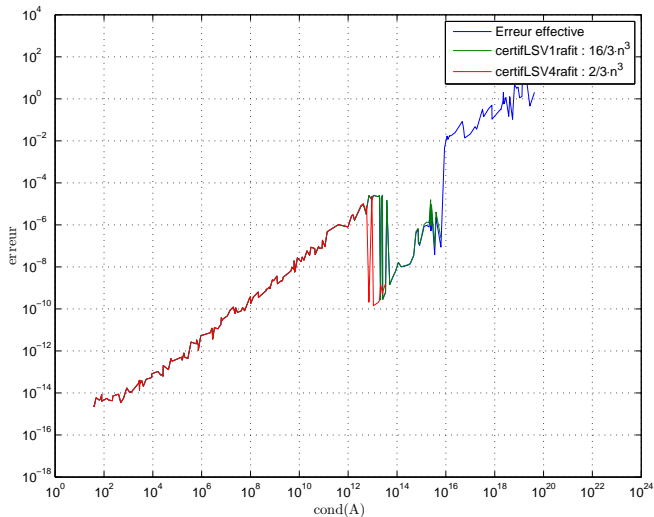
```
A = gallery('randsvd', n, 10^(k*rand));  
b = A*ones(n,1);
```

In the experiments, $n = 50$.


In the x-axis, we have $\kappa(A)$, in the y-axis we have the relative error. We use a logarithmic scale.

We change the tolerance τ which is the targeted relative error.

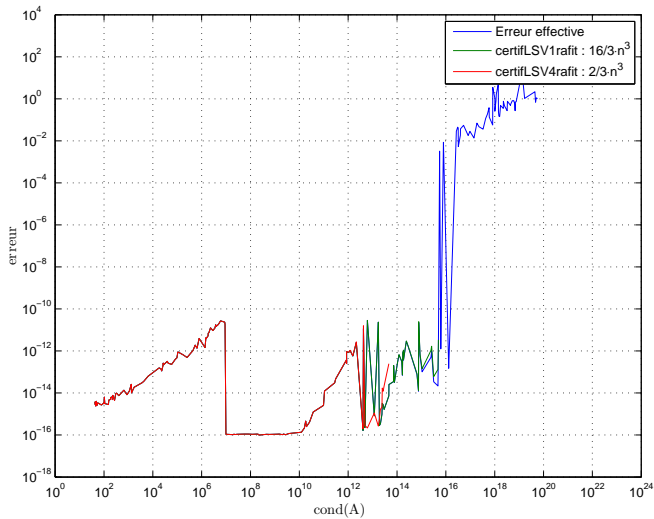
System 50×50 , $\text{tol} = 2^{-15} \approx 3 \cdot 10^{-5}$:




System 50×50 , $\text{tol} = 2^{-25} \approx 3 \cdot 10^{-8}$:



System 50×50 , $\text{tol} = 2^{-35} \approx 3 \cdot 10^{-11}$:



System 50×50 , $\text{tol} = 2^{-45} \approx 3 \cdot 10^{-14}$:



Summary

- The certified error bound behave like the true relative error for the computed solution: this behavior appears in increasing the accuracy of the computed solution by iterative refinement.
- Iterative refinement combined with verified algorithms make it possible to guarantee a “small” rigorous forward error as long as the condition number is “sufficiently small” compared to $1/u$. With the fastest algorithm we presented, it is still needed to double the number of floating-point operations.
- It would be interesting to provide efficient implementations of those algorithms as well as a deep experimental study in order to compare the performances with [DHK⁺06].

Outline

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

Outline

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

Statement of the problem

Let $A \in \mathbb{F}^{n \times n}$ be nonsingular. If we can compute an approximate inverse $R = \text{inv}(A)$ in floating-point arithmetic, and that $\|RA - I\| < 1$, we can show that A is nonsingular. It will be the case if $\kappa(A)$ is sufficiently small compared to u^{-1} .

The Kahan-Gastinel theorem states that

$$\kappa(A)^{-1} = \min \left\{ \frac{\|\Delta A\|}{\|A\|}, A + \Delta A \text{ is singular} \right\}.$$

If $\kappa(A) > u^{-1}$, a perturbation of A with normwise norm of order u can lead to a singular matrix...Moreover, if one wants to prove the nonsingularity of A , the strategy consisting in verifying that $\|RA - I\| < 1$ will likely to fail.

Statement of the problem

If $\kappa(A) > u^{-1}$, how can we prove the nonsingularity of A using floating-point arithmetic?

- We can still use multiprecision arithmetic
- There exists an algorithm from Rump [Rum09], in which it is only necessary to increase the precision used for matrix products.

Outline

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

Principle of the algorithm

Assume that we can perform some matrix operations exactly:

$R = \text{inv}(A)$	{approximate inverse, computed with work}
$P = RA$	{exact}
$X = P^{-1}$	{exact}
$R' = XR$	{exact}

In this case, $X = A^{-1}R^{-1}$, so $R' = XR = A^{-1}$.

In practice, we only authorize computation with finite precision so the previous algorithm cannot be used: but we can try to use some **multiplicative corrections** to compute a more accurate inverse of A .

Rump's algorithm

The idea of Rump is to locally use a precision greater than the working precision. Let $A, B \in \mathbb{F}^{n \times n}$ and $k \geq 2$.

The notation $P = fl_{k,1}(AB)$ means that the product AB is computed with precision u^k , then rounded to the working precision u :

$$\| fl_{k,1}(AB) - AB \| \leq u \| AB \| + nu^k \| A \| \| B \| + O(u^{k+1}).$$

The notation $\{P\} = fl_{k,k}(AB)$ means that AB is computed with precision u^k and the result is represented as a non-evaluated sum $\{P\} = \sum_{i=1}^k P_i$ ($P_i \in \mathbb{F}^{n \times n}$):

$$\| fl_{k,1}(AB) - AB \| \leq nu^k \| A \| \| B \| + O(u^{k+1}).$$

Moreover, if $\{P\} = \{P_1, \dots, P_\ell\}$,

$$fl_{k,1}(\{P\}B) = fl_{k,1}\left(\sum_{i=0}^{\ell} P_i B\right) \quad \text{and} \quad fl_{k,k}(\{P\}B) = fl_{k,k}\left(\sum_{i=0}^{\ell} P_i B\right).$$

Inversion of an ill-condition matrix

function $\{R^{(k)}\} = \text{InvIllCond}(A)$

$R^{(0)} = \text{fl}(\|A\|^{-1}) \cdot I$ {starting “approximate inverse”}

$k = 0$

repeat

$k = k + 1$

$P^{(k)} = \text{fl}_{k,1}(\{R^{(k-1)}\} \cdot A)$

$X^{(k)} = \text{inv}(P^{(k)})$

$\{R^{(k)}\} = \text{fl}_{k,k}(X^{(k)} \cdot \{R^{(k-1)}\})$

until $\text{cond}(P^{(k)}) < (100u)^{-1}$

If $\kappa(A) \gg u^{-1}$, Rump justifies in an heuristic way [Rum09] that

$$\text{cond}(\{R^{(k)}\}A) \approx u^{k-1} \text{cond}(A).$$

Moreover, we can hope that the algorithm terminates with
 $\|\{R^{(k)}\}A - I\| \leq 1/100$.

Let us take for example the following matrix A , such

$$\kappa(A_1) \approx 6.4 \cdot 10^{63}:$$

$$A = \begin{bmatrix} -5046135670319638 & -3871391041510136 & -5206336348183639 & -6745986988231149 \\ -640032173419322 & 8694411469684959 & -564323984386760 & -2807912511823001 \\ -16935782447203334 & -18752427538303772 & -8188807358110413 & -14820968618548534 \\ -1069537498856711 & -14079150289610606 & 7074216604373039 & 725796028397871 \end{bmatrix}.$$

For this matrix, we can observe:

k	$\text{cond}(R^{(k-1)})$	$\text{cond}(R^{(k-1)}A)$	$\text{cond}(P^{(k)})$	$\ I - R^{(k)}A\ $
2	$1.68 \cdot 10^{17}$	$2.73 \cdot 10^{49}$	$2.31 \cdot 10^{17}$	3.04
3	$1.96 \cdot 10^{32}$	$2.91 \cdot 10^{33}$	$2.14 \cdot 10^{17}$	5.01
4	$7.98 \cdot 10^{48}$	$1.10 \cdot 10^{17}$	$1.83 \cdot 10^{17}$	1.84
5	$6.42 \cdot 10^{64}$	8.93	8.93	$3.43 \cdot 10^{-16}$

- $\text{cond}(R^{(k-1)}A)$ decreases by a factor of order u at each iteration.
- By verifying $\|I - R^{(k)}A\| < 1$, one can show that A is nonsingular.

"Final version of the algorithm"

Inversion of an ill-conditioned matrix

function $\{R\} = \text{InvIllCond}(A)$

$\{R\} = \text{fl}(\|A\|^{-1}) \cdot I$; $P = X = \infty$; $k = 0$

repeat

finished = $(\|P\| \cdot \|X\| < (100u)^{-1})$

$k = k + 1$

$P = \text{fl}_{k,1}(\{R\} \cdot A)$

$X = \text{inv}(P)$

while "inversion failed" do

$P = P + \Delta P$ {random perturbations such that $|\Delta P| \leq u|P|$ }

$X = \text{inv}(P)$

done

$\{R\} = \text{fl}_{k,k}(X \cdot \{R\})$

until finished

Outline

- 1 Iterative refinement
 - Error analysis
 - Computation of a verified solution
 - Numerical experiments
- 2 Inversion of ill-conditioned matrices
 - Statement of the problem
 - Rump's algorithm
- 3 Conclusion

We presented four “fast” methods to compute a verified solution of linear systems with floating-point arithmetic.

- These methods require between 2 and 9 times for floating-point operations than the classic GE.
- All the methods can be implemented using the BLAS routines.
- They are efficient as long as the condition number is small compared to u^{-1} .

When the condition number is small compared to u^{-1} , if the accuracy of the computed solution is not sufficient, it is possible to use iterative refinement with mixed precision to obtain a normwise relative error of order u .

If the condition number is larger than u^{-1} , it is possible to use the Rump's inversion algorithm for preconditioning the matrix and obtain an approximate solution that can be verified.

References I



James Demmel, Yozo Hida, William Kahan, Xiaoye S. Li, Sonil Mukherjee, and E. Jason Riedy, Error bounds from extra-precise iterative refinement, *ACM Trans. Math. Softw.* **32** (2006), no. 2, 325–351.



Nicholas J. Higham, Accuracy and stability of numerical algorithms, second ed., Society for Industrial and Applied Mathematics (SIAM), Philadelphia, PA, 2002.



S. M. Rump, Inversion of extremely ill-conditioned matrices in floating-point, *Japan J. Indust. Appl. Math. (JJIAM)* **26** (2009), 249–277.